This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

DPON-385-E
March 2026

What is true, useful, and meaningful?

A leadership lens for AI

Sebastian Hafenbrädl

Introduction: Why leaders struggle to think clearly about AI

Artificial intelligence (AI) confronts leaders with a distinctive challenge. It is not simply a new tool but a class of systems that increasingly act, adapt, and participate in decisions once governed by human judgment alone. As a result, executives face a growing stream of consequential choices. They must decide which tools or copilots to roll out across the organization; which AI-supported processes, products, or ideas to pilot; which experiments to scale and which to abandon; and where to intervene top-down by redesigning processes, offerings, or entire areas around new AI capabilities. These decisions are recurring, interdependent, and often made under conditions of uncertainty and time pressure.

However, discussions surrounding these choices often feel urgent but unsatisfying. Technical assessments emphasize benchmark performance and rapid capability gains; business cases highlight productivity, scale, and competitive advantage; and ethical debates raise questions of accountability, professional judgment, and human agency. Rather than converging, these conversations frequently talk past one another. The difficulty is not a lack of intelligence, information, or intent. It is that leaders often try to answer different kinds of questions at the same time without fully recognizing how different those questions are.

In practice, decisions about AI tend to conflate three distinct conversations: what is true, what is useful, and what is meaningful. Each demands a different form of evidence, reasoning, and judgment. Each becomes distorted when answered through the lens of others.

This note is not about the latest developments in AI or about predicting which technologies will succeed. Instead, it addresses a more persistent challenge: how humans reason about transformative technologies that are fundamentally different from those they have encountered before. The aim is not to resolve all the difficulties of AI adoption but to remove one recurring obstacle to clarity—confusion between different kinds of questions—that often undermines individual judgment and, as a result, effective communication and alignment among decision-makers. When these distinctions are blurred, organizations do not merely make suboptimal decisions; they struggle to agree on what they are deciding in the first place.

This technical note was prepared by Professor Sebastian Hafenbrädl. March 2026.

Copyright © 2026 IESE. To order copies contact IESE Publishing via www.iesepublishing.com. Alternatively, write to publishing@iese.edu or call +34 932 536 558.
No part of this publication may be reproduced, stored in a retrieval system, made available on any LLM (i.e., ChatGPT), or transmitted in any form or by any means - electronic, mechanical, photocopying, recording, or otherwise - without the permission of IESE.

Last edited: 18/3/26

This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

What is true?

Technical capability, limits, and empirical reality

The first conversation leaders must hold concerns what is true about a technology. In the context of AI, this refers to the system's actual technical capabilities: what it can and cannot do, under what conditions, and with what degree of reliability.

In this note, "what is true" does not refer only to whether individual outputs are factually correct. It refers to the "empirical reality of the system"—how it actually behaves in practice. This includes its veracity and its characteristic failure modes, such as hallucination, bias, or overconfidence, as well as how humans tend to perceive and respond to those behaviors.

Technical capability is assessed through benchmarks, controlled evaluations, and related forms of evidence. These measures are essential, but they are also conditional. They reflect specific tasks, datasets, and evaluation regimes (for which models might be optimized), and they capture performance under observed conditions at a given point in time. High scores signal progress, not comprehensive understanding.

In practice, AI capabilities are uneven. Systems may perform extremely well on some tasks while failing unexpectedly on others—a pattern often described as a jagged frontier of performance. As a result, answering what is true is rarely a matter of declaring that a system works or does not work. It requires judgment about *where* it works, *how reliably*, and *with what consequences when it fails*.

A further challenge for leaders is choosing the right flight altitude at which to engage technical information. Immersion in technical news—model releases, marginal benchmark gains, or incremental performance improvements—quickly becomes a form of AI gossip. There are more developments than any executive can productively track, and doing so rarely improves decision quality.

What matters instead is conceptual literacy: understanding classes of capability, typical failure modes, limits to generalization, and how performance changes when systems are embedded in organizational contexts. This level of understanding is more valuable than familiarity with the latest metrics or decimal point improvements.

Human responses are part of this empirical reality. Studies show, for example, that people may perceive AI-generated faces as more humanlike than real photographs or evaluate AI-generated medical responses as more empathic than those written by physicians. Such findings matter because perceptions shape trust, reliance, and action. At the same time, these studies isolate specific mechanisms under defined conditions. Organizational settings are multidetermined, combining context, incentives, accountability, and social interaction. Evidence about human responses is therefore informative but never self-interpreting. Part of the truth about an AI system may be that it sometimes produces confident falsehoods—and that humans tend to trust them anyway.

Taken together, the question of what is true establishes a necessary but limited foundation. It describes capabilities and responses as they are currently observed, at one point in time, in a given context—not outcomes, as they will unfold over time. Treating these assessments as conclusions rather than inputs risks confusing evidence with foresight.

This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

What is useful?

Organizational value, dynamics, and strategy

The second conversation leaders must hold concerns what is useful about a technology. In the context of AI, this refers not to technical capability itself but to the ways in which that capability is translated into personal, organizational, economic, and strategic value.

Unlike the question of truth, usefulness is inherently forward-looking. It asks how a technology will shape outcomes over time: how it will change workflows, decision processes, cost structures, coordination, and competitive dynamics. Answers to this question are not discovered in benchmarks but are constructed through assumptions about behavior, incentives, and adaptation.

A common mistake is to treat usefulness as a direct function of technical performance. Improvements in accuracy, speed, or scale are assumed to translate automatically into productivity gains or strategic advantage. In practice, this translation is neither automatic nor linear. Value depends on how technologies are embedded in processes, how roles and responsibilities shift, and how people adjust their behavior in response.

Many AI initiatives disappoint not because the technology fails but because organizations underestimate second-order effects. Designing for usefulness is akin to designing a building; choices are made with intended use in mind, yet once in place, the structure itself begins to shape behavior. Tools meant to improve efficiency can alter accountability, change risk-taking, or redirect attention—effects that often become visible only after deployment.

Therefore, usefulness is not synonymous with efficiency. A technology is useful only insofar as it advances an organization's strategy. If AI makes everyone 10% faster at writing emails, it does not necessarily create an advantage—it might establish a new baseline. When technologies reduce the cost of commodity activities, they change the floor, not the ceiling.

Usefulness is also dynamic. Early gains may erode as competitors adopt similar tools, as systems require increasing governance, or as employees learn to work around—or rely excessively on—automated outputs. Conversely, modest initial benefits may compound as organizations redesign processes, build complementary capabilities, and learn how to use the technology more effectively. Judgments about usefulness, therefore, cannot be reduced to a single business case or pilot result.

Another recurring trap is treating usefulness as value in isolation. Gains in one area may impose burdens elsewhere, and benefits to one group may create costs for another. The overall value might be bottlenecked by the least efficient process step. Leaders must therefore attend not only to aggregate value creation but also to how value and burden are distributed across roles, functions, and time.

To engage the question of usefulness productively, leaders should ask not whether a technology promises value but "how that value is expected to emerge, for whom, and under what conditions". This includes attention to incentives, learning curves, organizational redesign, and the possibility that anticipated benefits may change as the system and its context evolve.

Crucially, answering what is useful does not determine whether a technology should be pursued. Usefulness speaks to expected outcomes, not legitimacy or desirability. When projections of value are treated as sufficient justification, leaders risk confusing plausible futures with acceptable ones.

This document is an authorized copy for the course Master in Management 2027 - Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

What is meaningful?

Legitimacy, human meaning, and social judgment

The third conversation leaders must hold concerns what is meaningful about a technology. In the context of AI, this refers to how its use is evaluated socially and morally—whether it is seen as appropriate, legitimate, or aligned with shared understandings of human purpose and responsibility.

As with the question of truth, meaningfulness concerns technology itself and "human nature in relation to it". Leaders shape technologies through design and deployment choices. In turn, the sustained use of these technologies shapes how people work, decide, relate to one another, and understand their own roles. Meaning arises not only from outcomes but also from participation in these evolving practices.

A critical distinction is required—one that is often blurred in discussions about AI. Meaningfulness can stem from the experience of "working with or alongside a technology", and it can also stem from "the benefits that work provides to others". These two sources of meaning do not always align.

Consider an idealized cancer researcher who finds purpose primarily in contributing to cures that alleviate human suffering. From this perspective, the researcher might welcome AI systems that automate large portions of their work if doing so accelerates discovery and broadens access to treatment. Here, meaning is located less in the activity of work itself than in its contribution to a broader human good.

In contrast, other forms of work derive meaning from the exercise of judgment, care, skill, or professional identity. In these cases, displacement or reconfiguration of human involvement may be experienced not as progress but as erosion. Whether AI is seen as augmenting or undermining meaning depends not only on what it achieves but on how it redistributes agency, responsibility, and recognition.

These tensions become most visible at the margins—where AI systems approach and, in some domains, exceed human performance. Even if a future in which AI can perform all human jobs better and cheaper remains speculative or unreachable, leaders increasingly confront local versions of this question—cases in which automation promises clear benefits but unsettles deeply held assumptions about the value of human contribution.

Meaningfulness is therefore not resolved by asking whether work remains, but by asking "what kind of participation remains meaningful", and why. These judgments are collective rather than individual. They depend on professional norms, organizational cultures, and societal expectations about dignity, contribution, and responsibility.

Importantly, meaningfulness is forward-looking but not predictive. It concerns the kind of world that repeated deployment choices make more likely—not only in outcomes but also in who humans become through sustained interaction with intelligent systems. There are no benchmarks for such judgments and no business cases that can settle them.

Crucially, answering what is meaningful neither overrides nor concludes conversations about truth and usefulness. It is a distinct form of judgment that must be held alongside them. When questions of meaning are ignored, deferred, or reduced to compliance, leaders risk discovering that technologies that are true and useful nonetheless fail to command legitimacy—or to sustain the human commitments on which organizations ultimately depend.

This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

When the questions are not kept distinct

When the questions of what is true, what is useful, and what is meaningful are not kept distinct, predictable errors in judgment arise. These failures rarely reflect poor intent. They stem from allowing answers to one kind of question to substitute for another. In practice, these questions often function as informal gates. Leaders frequently begin with usefulness—asking what value a technology might create—and then work backward to justify assumptions about truth or meaning. This sequencing is risky. Judgments about usefulness are least reliable when they are made before there is sufficient clarity about empirical reality or about what is morally desirable and can be justified to those affected.

When assessments of truth are colored by assumptions about usefulness or meaning, leaders risk wishful thinking. Evidence about technical or behavioral capabilities may be discounted because the implications feel undesirable. For example, if leaders believe that empathic responses from AI systems are unlikely to benefit patients or that morally patients ought to interact with human physicians, they may conclude that AI cannot provide empathic responses—despite empirical evidence suggesting otherwise. In such cases, normative or instrumental preferences shape beliefs about what is true, undermining the discipline required to follow the evidence where it leads.

Conversely, when judgments about usefulness are treated as direct extensions of technical capability, leaders overlook the dynamic nature of value creation. What appears as a competitive advantage at one point in time may quickly become a commodity. The usefulness of AI often depends on ecosystem dynamics, network effects, and organizational learning that unfold over time. Moreover, usefulness is not intrinsic to the technology; it must stem from strategy—from the organization's mission, value proposition, and choices about where and how to compete. Absent this anchoring, projections of value risk are technically sophisticated but strategically ungrounded.

Finally, when usefulness is not separated from meaningfulness, economic logic can quietly displace moral judgment. AI may lower the cost of certain forms of intelligence while raising the value of complementary human capabilities. It may enable efficiency gains that are compelling in cost–benefit terms. However, organizations are not merely economic instruments. They are moral actors with purposes that cannot be reduced to optimization. When usefulness is allowed to stand in for meaning, leaders risk building organizations that are efficient but misaligned with the very purposes that justify their existence.

In each case, the failure is not in asking the wrong questions but in answering the right questions in the wrong order—or allowing one kind of answer to do the work of another. Keeping these conversations distinct does not eliminate trade-offs. It makes them visible, discussable, and subject to judgment rather than drift.

This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

Conclusion: Judgment in an age of powerful tools

AI confronts leaders not with a single decision but with a persistent challenge of judgment. That challenge cannot be met by collapsing questions of truth, usefulness, and meaningfulness into one conversation. Each addresses a different dimension of leadership responsibility, each requires a different form of reasoning, and each becomes misleading when it substitutes for the others.

A helpful way to think about this challenge is through the metaphor of building a house. Leaders make design choices based on intended use—how many rooms are needed, how spaces should connect, and what activities they should support. However, once built, the house begins to shape behavior in return. Small design decisions influence how people move, gather, and live together. Over time, inhabitants discover whether the house is merely functional or whether it is a place they want to inhabit and defend.

Technologies such as AI are similar. Assessments of what is true correspond to understanding the materials and structural limits: what the system can bear, where it holds, and where it fails. Judgments about what is useful resemble decisions about layout and use: how the technology is integrated into workflows, how it changes behavior over time, and whether it supports the organization's strategy and mission. Questions of what is meaningful ask something deeper: whether the resulting practices align with shared values, justify themselves to those affected, and contribute to a form of work and organization worth sustaining.

The difficulty is that the pace and power of contemporary technologies increasingly outstrip our capacity to imagine their consequences. Philosopher Günther Anders described this as the "Promethean Gap": the widening distance between what humans can create and what they can fully comprehend, emotionally and morally. In the context of AI, leaders can authorize deployments whose effects on work, agency, and responsibility become visible only over time—often after systems are scaled and routines have hardened.

The purpose of the lens offered here is not to close this gap by prediction or control. It is to narrow it through disciplined inquiry. By keeping questions of truth, usefulness, and meaningfulness distinct—and by holding them in tension rather than resolving them prematurely—leaders create the conditions for more responsible judgment. They slow down where evidence is required, look ahead where dynamics matter, and pause where legitimacy and purpose are at stake.

In the end, leadership in the age of AI is less about choosing the right answers than about asking the right questions—in the right order and at the right level of abstraction. Technologies may increasingly shape how organizations function. However, whether those organizations remain effective, legitimate, and human is still a matter of judgment—and that responsibility cannot be delegated.

This document is an authorized copy for the course Master in Management 2027 - Organizational Behavior - B taught by prof. Villamor, Isabel at IESE B.S.

Exhibit

Executive diagnostic: Anchoring the three conversations

When evaluating an AI initiative, executive discussions often drift—toward technical enthusiasm, abstract concern, or premature conclusions. The questions below are intended to "anchor the conversation" by keeping the three forms of judgment—truth, usefulness, and meaningfulness—distinct.

These questions do not comprise a checklist. They are prompts for disciplined inquiry.

1. Auditing what is true

The reality check

- What evidence do we have about this system's performance "under our conditions of use" rather than under idealized or benchmark settings?
- How does the system tend to fail—and are those failures visible, containable, and understandable?
- Where does this system perform reliably, and where does performance become uneven or unpredictable?
- What do we know about "human responses" to this system (e.g., trust, reliance, and overconfidence), and how might those responses shape outcomes?

2. Auditing what is useful

The strategic check

- How is value expected to emerge over time—not just at initial deployment?
- If similar capabilities become widely available, what, if anything, sustains our advantage?
- How might this technology shift burdens and benefits across roles, teams, or career stages?
- If reliance on this system increases, which human capabilities or organizational checks might weaken—and which might become more valuable?
- How does this initiative reinforce, rather than dilute, our strategy, mission, and value proposition?

3. Auditing what is meaningful

The legitimacy check

- Does this use of AI change "what gives the work meaning" to those involved—and in what way?
- Does it automate tasks that obstruct human contributions or judgments that define professional responsibility?
- How does this deployment affect accountability: Who explains outcomes and who owns consequences?
- Could we plausibly explain and defend this use of AI to those affected by it?
- What kind of organization does repeated reliance on this technology make over time?

A final discipline

Have these questions been discussed "separately" before being weighed together?